

LangChain is an orchestration framework for developing applications powered by large language models. It provides tools for chaining together different components like models, APIs, and data sources. RAG (Retrieval-Augmented Generation) is a technique that connects an LLM to external knowledge bases. This allows the model to retrieve relevant information from documents like PDFs or web pages before generating an answer, which helps prevent hallucinations.